


CS 4495 Computer Vision

RANdom **SA**mples **C**onsensus

Aaron Bobick
School of Interactive
Computing




Administrivia

- PS 3:
 - Check Piazza - learn about least squares solutions and pseudo inverses
 - For 1.3, instead of using 4, 8 and 16 points and looking at the average residuals of 5 points, use 4, 8, and 15 – since there are only 20 points!!
 - Due TUES October 18. Don't blow it!
- PS 4 will be out the 20th and due the 30th.


Matching with Features

- Detect feature points in both images



An introductory example:

Harris corner detector



C.Harris, M.Stephens. “A Combined Corner and Edge Detector”. 1988

Harris Detector: Mathematics

$$M = A^T A = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix}$$

Measure of corner response:

$$R = \det M - k (\text{trace } M)^2$$


$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

(k – empirical constant, $k = 0.04$ - 0.06)


Harris Detector: Mathematics

- R depends only on eigenvalues of M
- R is large for a **corner**
- R is negative with large magnitude for an **edge**
- $|R|$ is small for a **flat** region




Key point localization

- General idea: find robust extremum (maximum or minimum) both in space and in scale.
- SIFT specific suggestion: use DoG pyramid to find maximum values (remember edge detection?) – then eliminate “edges” and pick only corners.
- More recent: use Harris detector to find maximums in space and then look at the Laplacian pyramid (we’ll do this later) for maximum in scale.




Each point is compared to its 8 neighbors in the current image and 9 neighbors each in the scales above and below.


Remove low contrast, edge bound



Extrema points




Contrast $> C$




Not only on edge

Scale Invariant Detectors

- **Harris-Laplacian**¹
Find local maximum of:
 - Harris corner detector in space (image coordinates)
 - Laplacian in scale



- **SIFT (Lowe)**²
Find local maximum of:
 - Difference of Gaussians in space and scale




¹ K.Mikolajczyk, C.Schmid. "Indexing Based on Scale Invariant Interest Points". ICCV 2001

² D.Lowe. "Distinctive Image Features from Scale-Invariant Keypoints". Accepted to IJCV 2004

Point Descriptors

- We know how to detect points
- Next question: How to match them?




Point descriptor should be:

1. Invariant
2. Distinctive

Another version of the problem...

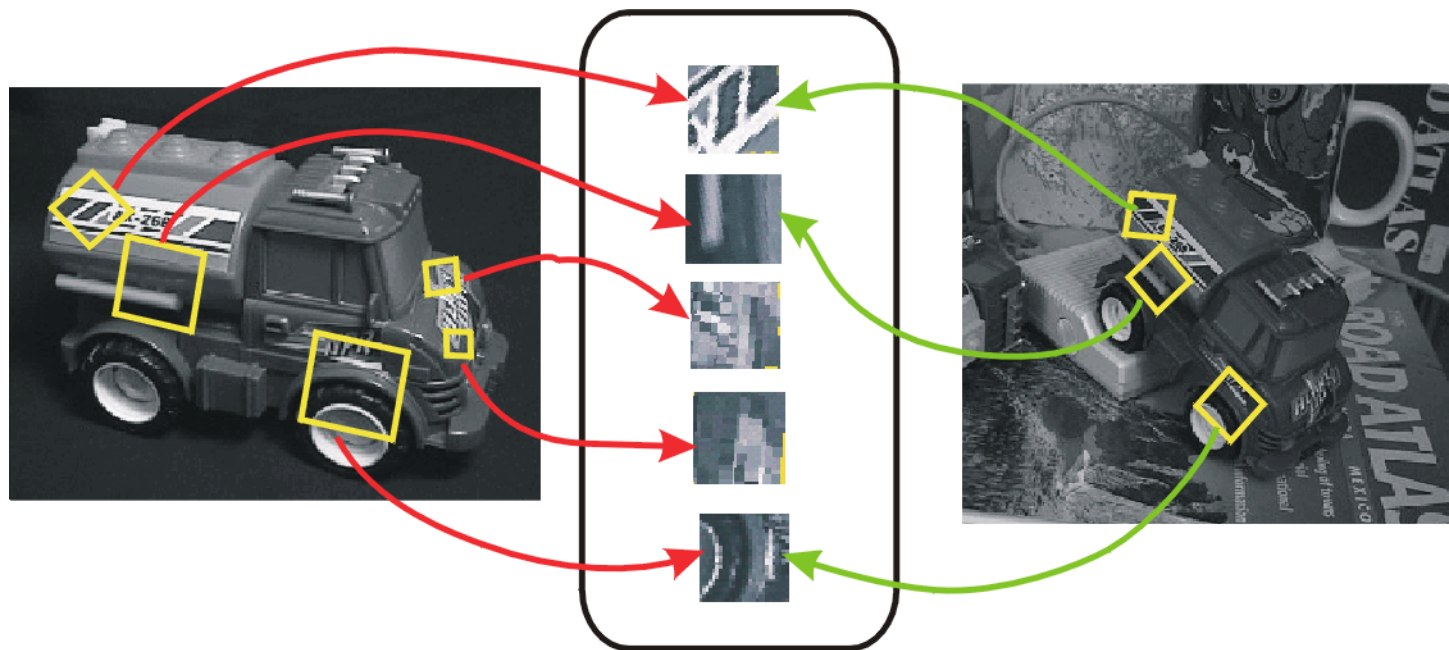
Want to find

... in here



Idea of SIFT

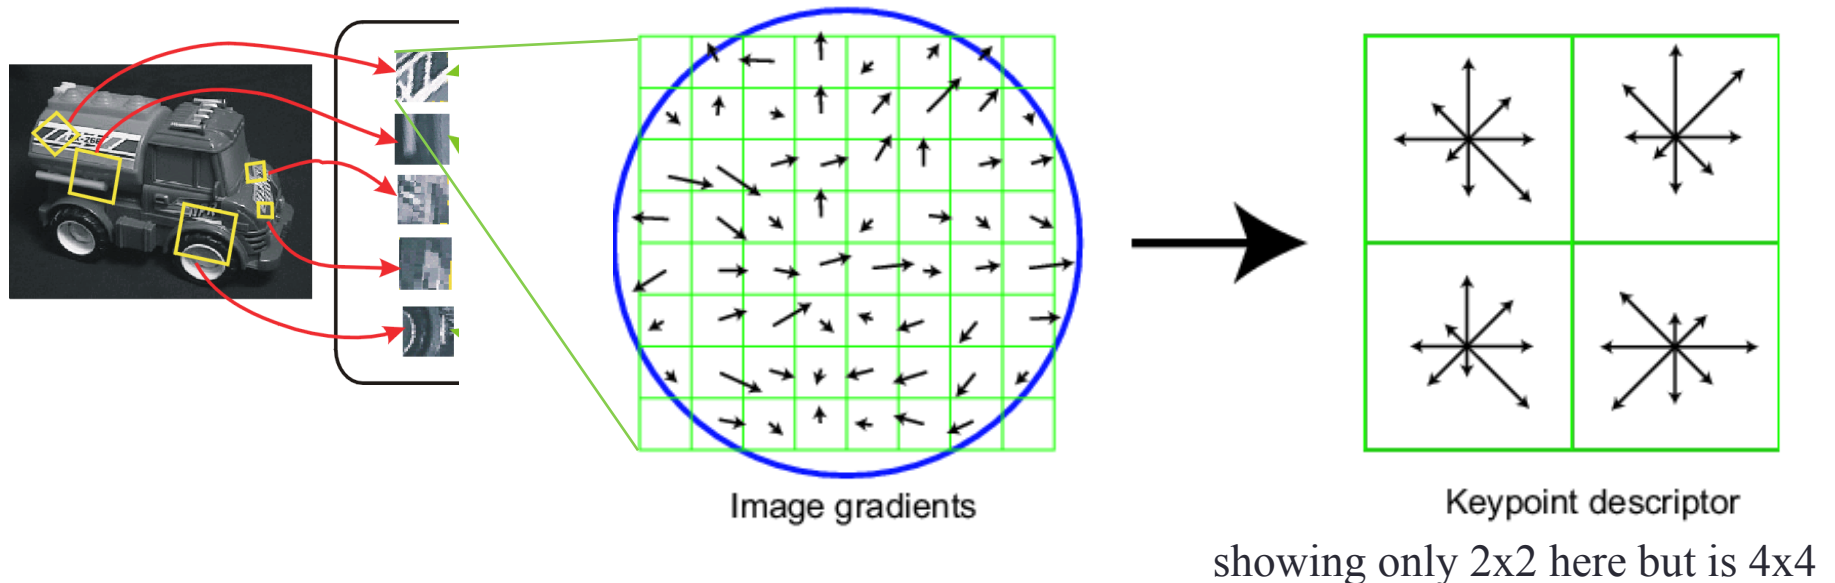
- Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, and other imaging parameters





SIFT Features

SIFT vector formation

- 4x4 array of gradient orientation histograms over 4x4 pixels
 - not really histogram, weighted by magnitude
- 8 orientations x 4x4 array = 128 dimensions
- Motivation: some sensitivity to spatial layout, but not too much.




Feature-based alignment outline




- Extract features

Feature-based alignment outline




- Extract features
- Compute *putative* matches

Feature-based alignment outline



- Extract features
- Compute **putative** matches
- Loop:
 - Hypothesize transformation T

Feature-based alignment outline



- Extract features
- Compute **putative** matches
- Loop:
 - Hypothesize transformation T
 - Verify transformation (search for other matches consistent with T)

Feature-based alignment outline



- Extract features
- Compute *putative* matches
- Loop:
 - *Hypothesize* transformation T
 - *Verify* transformation (search for other matches consistent with T)


How to get “putative” matches?

Feature matching

- Exhaustive search
 - for each feature in one image, look at *all* the other features in the other image(s) – pick best one
- Hashing
 - compute a short descriptor from each feature vector, or hash longer descriptors (randomly)
- Nearest neighbor techniques
 - k -trees and their variants


Feature-space outlier rejection

- Let's not match all features, but only these that have “similar enough” matches?
- How can we do it?
 - $\text{SSD}(\text{patch1}, \text{patch2}) < \text{threshold}$
 - How to set threshold?



Feature-space outlier rejection

- A better way [Lowe, 1999]:
 - 1-NN: SSD of the closest match
 - 2-NN: SSD of the second-closest match
 - Look at how much better 1-NN is than 2-NN, e.g. 1-NN/2-NN
 - That is, is our best match so much better than the rest?



Feature matching


- Exhaustive search
 - for each feature in one image, look at *all* the other features in the other image(s) – pick best one
- Hashing
 - compute a short descriptor from each feature vector, or hash longer descriptors (randomly)
- Nearest neighbor techniques
 - k -trees and their variants
- *Problem: Even when pick best match, still lots (and lots) of wrong matches – “outliers”*

Another way to remove mistakes

- Why are we doing matching?
 - To compute a model of the relation between entities
- So this is really “model fitting”

Fitting

- Choose a parametric model to represent a set of features – *remember this???*



simple model: lines



simple model: circles



complicated model: car

Fitting: Issues

Case study: Line detection




- **Noise** in the measured feature locations
- **Extraneous data:** clutter (outliers), multiple lines
- **Missing data:** occlusions

Least squares line fitting

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$

- Data: $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation: $y_i = mx_i + b$
- Find (m, b) to minimize



$$E = \sum_{i=1}^n \left(y_i - \begin{bmatrix} x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right)^2 = \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right\|^2 = \|Y - XB\|^2$$

$$= (Y - XB)^T (Y - XB) = Y^T Y - 2(XB)^T Y + (XB)^T (XB)$$

$$\frac{dE}{dB} = 2X^T XB - 2X^T Y = 0$$

$$X^T XB = X^T Y$$


Normal equations: least squares solution to
 $XB=Y$

Problem with “vertical” least squares

- Not rotation-invariant
- Fails completely for vertical lines


Total least squares

- Distance between point (x_i, y_i) and line $ax+by=d$ ($a^2+b^2=1$): $|ax_i + by_i - d|$



Total least squares

- Distance between point (x_i, y_i) and line $ax+by=d$
- Find (a, b, d) to minimize the sum of squared perpendicular distances



$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$

Total least squares

- Distance between point (x_i, y_i) and line $ax+by=d$
- Find (a, b, d) to minimize the sum of squared perpendicular distances


$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$

$$\frac{\partial E}{\partial d} = \sum_{i=1}^n -2(ax_i + by_i - d) = 0$$

$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = (UN)^T (UN)$$

$$\frac{dE}{dN} = 2(U^T U)N = 0$$

Solution to $(U^T U)N = 0$, subject to $\|N\|^2 = 1$: eigenvector of $U^T U$ associated with the smallest eigenvalue (least squares solution to *homogeneous linear system* $UN = 0$)



$$d = \frac{a}{n} \sum_{i=1}^n x_i + \frac{b}{n} \sum_{i=1}^n y_i = a\bar{x} + b\bar{y}$$

Total least squares


$$U = \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix}$$

$$U^T U = \begin{bmatrix} \sum_{i=1}^n (x_i - \bar{x})^2 & \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) \\ \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) & \sum_{i=1}^n (y_i - \bar{y})^2 \end{bmatrix}$$

second moment matrix


Least squares: Robustness to noise

- Least squares fit to the red points:



Least squares: Robustness to noise

- Least squares fit with an outlier:




Problem: squared error heavily penalizes outliers

Robust estimators

- General approach: minimize $\sum_i \rho(r_i(x_i, \theta); \sigma)$


$r_i(x_i, \theta)$ – residual of i th point w.r.t. model parameters θ

ρ – robust function with scale parameter σ




The robust function ρ behaves like squared distance for small values of the residual u but saturates for larger values of u

Choosing the scale: Just right




The effect of the outlier is minimized

Choosing the scale: Too small



The error value is almost the same for every point and the fit is very poor

Choosing the scale: Too large




Behaves much the same as least squares

“Find consistent matches”???

- Some points (many points) are static in the world
- Some are not
- Need to find the right ones so can compute pose.
- Well tried approach:
 - Random Sample Consensus (RANSAC)

Simpler Example

- Fitting a straight line



Discard Outliers


- No point with $d > t$
- RANSAC:
 - RANdom SAmple Consensus
 - Fischler & Bolles 1981
 - Copes with a large proportion of outliers

M. A. Fischler, R. C. Bolles. [Random Sample Consensus: A Paradigm for Model Fitting with Applications to Image Analysis and Automated Cartography](#). Comm. of the ACM, Vol 24, pp 381-395, 1981.

RANSAC

(**RAN**dom **SA**mples **C**onsensus) :

Fischler & Bolles in '81.




Algorithm:

1. **Sample** (randomly) the number of points required to fit the model
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

RANSAC

Line fitting example




Algorithm:

1. **Sample** (randomly) the number of points required to fit the model ($\#=2$)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

RANSAC

Line fitting example



Algorithm:


1. **Sample** (randomly) the number of points required to fit the model ($\#=2$)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

RANSAC

Line fitting example

$$N_I = 6$$




Algorithm:

1. **Sample** (randomly) the number of points required to fit the model ($\# = 2$)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

RANSAC



Algorithm:

1. **Sample** (randomly) the number of points required to fit the model ($\# = 2$)
2. **Solve** for model parameters using samples
3. **Score** by the fraction of inliers within a preset threshold of the model

Repeat 1-3 until the best model is found with high confidence

Best Line has most support

- More support -> better fit

In General

- Fit a more general model
- Sample = minimal subset - s
 - Translation: pick one point pair
 - Homography (for plane) – pick 4 point pairs
 - Fundamental matrix – pick 7 point pairs
- Algorithm
 - Randomly select s points
 - Instantiate a model
 - Get consensus set S_i
 - If $|S_i| > T$, terminate and return model
 - Repeat for N trials, return model with $\max |S_i|$

Distance Threshold

- Requires noise distribution
- Location: Gaussian noise with σ
- Distance: Chi-squared distribution with DOF m
 - 95% cumulative:
 - Line, F: $m=1$, $t=3.84 \sigma^2$
- I.e. \rightarrow 95% prob that $d < t$ when point is inlier

How many samples ?

- We want: at least one sample with all inliers
- Can't guarantee: probability p
- E.g. $p = 0.99$

Choosing the parameters

- Initial number of points s
 - Typically minimum number needed to fit the model
- Distance threshold t
 - Choose t so probability for inlier is p (e.g. 0.95)
 - Zero-mean Gaussian noise with std. dev. σ : $t^2 = 3.84\sigma^2$
- Number of samples N
 - Choose N so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$) (outlier ratio: e)

Calculate N

- s – number of points to compute solution
- p – probability of success
- e – proportion outliers, so % inliers = $(1-e)$
- $P(\text{sample set with all inliers}) = (1-e)^s$
- $P(\text{sample set will have at least one outlier}) = 1 - (1-e)^s$
- $P(\text{all } N \text{ samples have outlier}) = (1 - (1-e)^s)^N$
- We want $P(N \text{ samples an outlier}) < 1-p$
- $(1 - (1-e)^s)^N < 1-p$

$$N > \log(1-p) / \log(1 - (1-e)^s)$$

Samples required for inliers only in a sample


- $P=0.99$

- $s=2, \varepsilon=5\%$ $\Rightarrow N=2$
- $s=2, \varepsilon=50\%$ $\Rightarrow N=17$
- $s=4, \varepsilon=5\%$ $\Rightarrow N=3$
- $s=4, \varepsilon=50\%$ $\Rightarrow N=72$
- $s=8, \varepsilon=5\%$ $\Rightarrow N=5$
- $s=8, \varepsilon=50\%$ $\Rightarrow N=1177$


s	proportion of outliers e						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177

- $N = f(\varepsilon)$, *not the number of points*
- N increases steeply with s


$$N > \log(1 - p) / \log(1 - (1 - e)^s)$$




Matching features



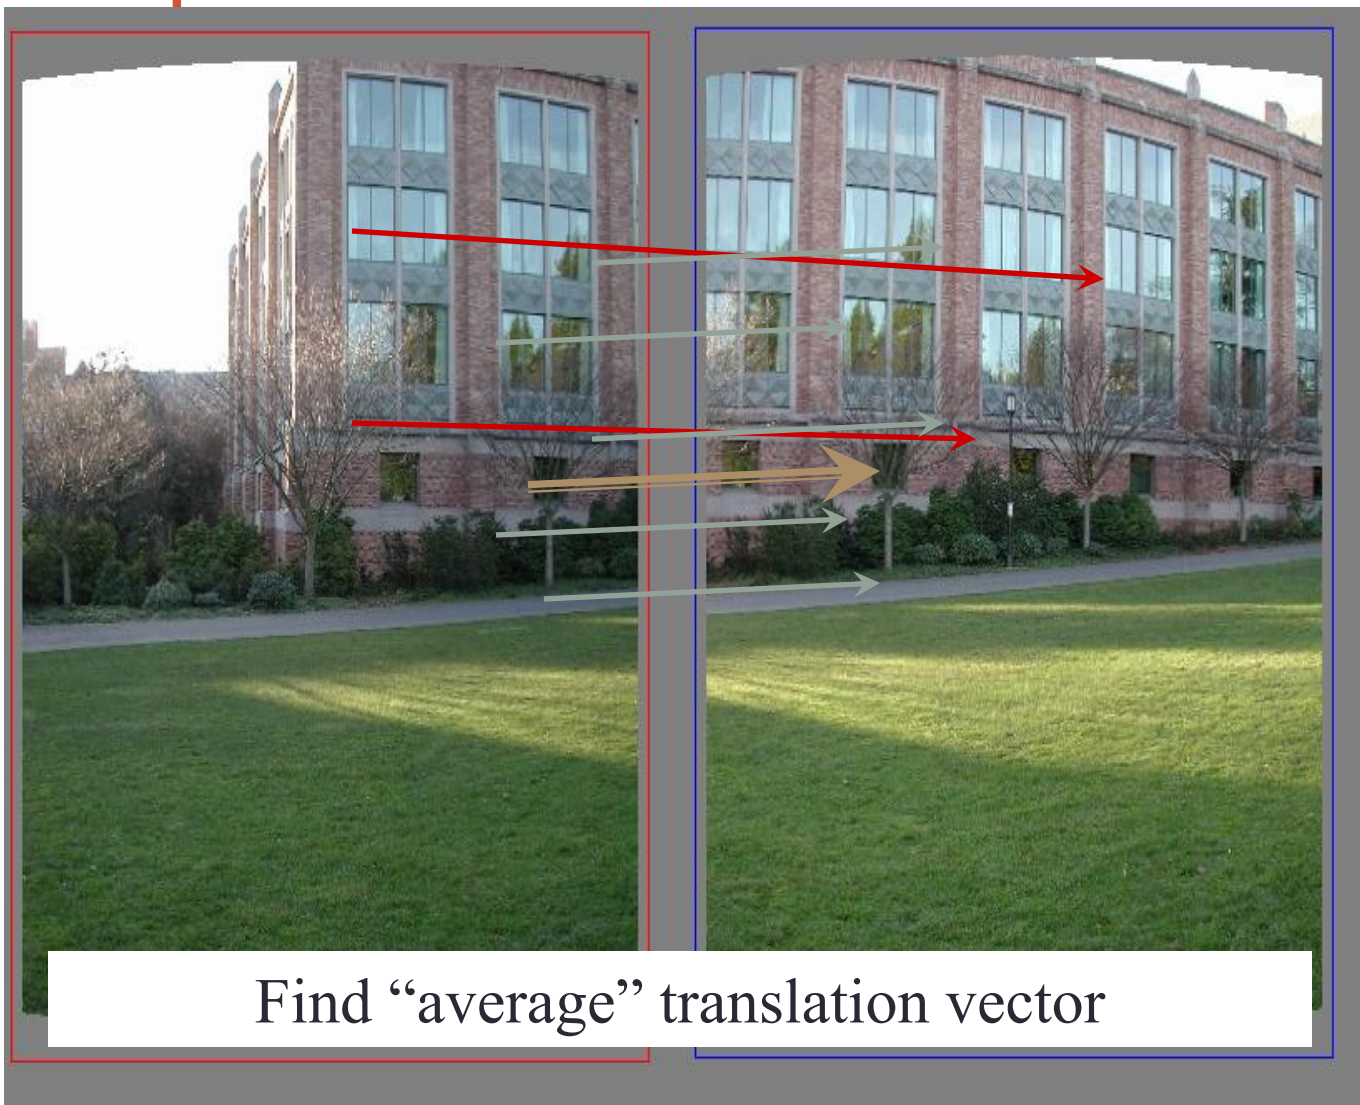
Random Sample Consensus



Random Sample Consensus



Least squares fit




RANSAC for estimating homography


- RANSAC loop:
 1. Select four feature pairs (at random)
 2. Compute homography H (exact)
 3. Compute *inliers* where $SSD(p_i', H p_i) < \varepsilon$
 4. Keep largest set of inliers
 5. Re-compute least-squares H estimate on all of the inliers

2D transformation models


- Similarity
(translation, scale, rotation)



- Affine



- Projective
(homography)



Adaptively determining the number of samples

- Inlier ratio e is often unknown a priori, so pick worst case, e.g. 50%, and adapt if more inliers are found, e.g. 80% would yield $e=0.2$
- Adaptive procedure:
 - $N=\infty$, sample_count = 0
 - While $N > \text{sample_count}$
 - Choose a sample and count the number of inliers
 - Set $e = 1 - (\text{number of inliers})/(\text{total number of points})$
 - Recompute N from e :

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$

- Increment the sample_count by 1

RANSAC conclusions

Good

- Simple and general
- Applicable to many different problems, often works well in practice
- Robust to outliers
- Applicable for larger number of parameters than Hough transform
- Parameters are easier to choose than Hough transform

Bad

- Computational time grows quickly with fraction of outliers and number of parameters
- Not good for getting multiple fits

Common applications

- Computing a homography (e.g., image stitching)
- Estimating fundamental matrix (relating two views)
- Every problem in robot vision