

# CS 4495 Computer Vision

## *Motion Models*

Aaron Bobick  
School of Interactive  
Computing

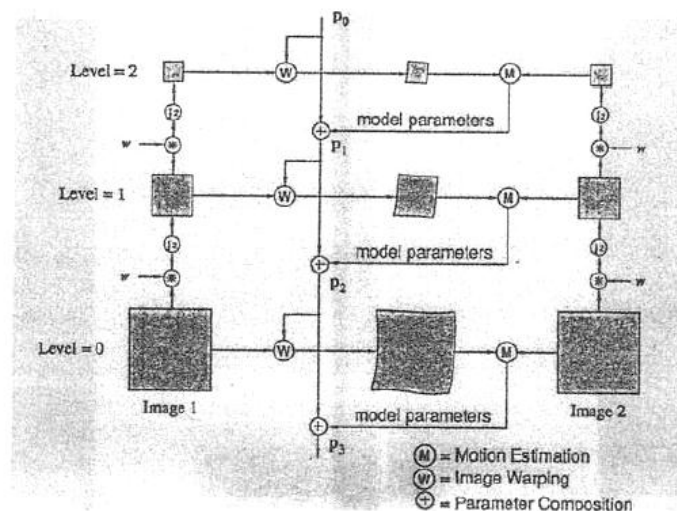


Fig. 1. Diagram of the hierarchical motion estimation framework.

# Outline



- 2D Motion models
  - Bergen, '92
  - Pyramids
  - Layers
- Motion fields from 3D motions
- Parametric motion

# Last time: Arbitrary visual motion





Many slides adapted from S. Seitz, R. Szeliski, M. Pollefeys, K. Grauman and others...

# Motion estimation: Optical flow



Will start by estimating motion of each pixel separately  
Then will consider motion of entire image

# Problem definition: optical flow


 $I(x, y, t)$ 

 $I(x, y, t + 1)$ 

How to estimate pixel motion from image  $I(x, y, t)$  to  $I(x, y, t + 1)$  ?

- Solve pixel correspondence problem
  - given a pixel in  $I(x, y, t)$ , look for **nearby** pixels of the **same color** in  $I(x, y, t + 1)$

## Key assumptions

- **color constancy**: a point in  $I(x, y, t)$  looks the same in  $I(x, y, t + 1)$ 
  - For grayscale images, this is brightness constancy
- **small motion**: points do not move very far

This is called the optical flow problem

# Optical flow equation

- Combining these two equations

$$\begin{aligned} 0 &= I(x+u, y+v, t+1) - I(x, y, t) && \text{shorthand: } I_x = \frac{\partial I}{\partial x} \\ &\approx I(x, y, t+1) + I_x u + I_y v - I(x, y, t) \\ &\approx [I(x, y, t+1) - I(x, y, t)] + I_x u + I_y v \\ &\approx I_t + I_x u + I_y v \\ &\approx I_t + \nabla I \cdot \langle u, v \rangle \end{aligned}$$

In the limit as  $u$  and  $v$  go to zero, this becomes exact

$$0 = I_t + \nabla I \cdot \langle u, v \rangle$$

*Brightness constancy constraint equation*

$$I_x u + I_y v + I_t = 0$$

# Optical flow equation


$$0 = I_t + \nabla I \cdot \langle u, v \rangle \quad \text{or} \quad I_x u + I_y v + I_t = 0$$

- Q: how many unknowns and equations per pixel?


**2 unknowns, one equation**

Intuitively, what does this constraint mean?

- The component of the flow in the gradient direction is determined
- The component of the flow parallel to an edge is unknown



# Aperture problem





# Solving the aperture problem

- How to get more equations for a pixel?
  - Basic idea: impose additional constraints
    - most common is to assume that the flow field is smooth locally
    - one method: pretend the pixel's neighbors have the same (u,v)
      - If we use a 5x5 window, that gives us 25 equations per pixel!

$$0 = I_t(\mathbf{p}_i) + \nabla I(\mathbf{p}_i) \cdot [u \ v]$$

$$\begin{array}{c}
 \left[ \begin{array}{cc} I_x(\mathbf{p}_1) & I_y(\mathbf{p}_1) \\ I_x(\mathbf{p}_2) & I_y(\mathbf{p}_2) \\ \vdots & \vdots \\ I_x(\mathbf{p}_{25}) & I_y(\mathbf{p}_{25}) \end{array} \right] \begin{bmatrix} u \\ v \end{bmatrix} = - \left[ \begin{array}{c} I_t(\mathbf{p}_1) \\ I_t(\mathbf{p}_2) \\ \vdots \\ I_t(\mathbf{p}_{25}) \end{array} \right] \\
 \begin{array}{ccc} \mathbf{A} & \mathbf{d} & \mathbf{b} \\ 25 \times 2 & 2 \times 1 & 25 \times 1 \end{array}
 \end{array}$$

# Lukas-Kanade flow

- Prob: we have more equations than unknowns

$$\begin{array}{ccc} A & d = b & \longrightarrow \text{minimize } \|Ad - b\|^2 \\ 25 \times 2 & 2 \times 1 & 25 \times 1 \end{array}$$

Solution: solve least squares problem

- minimum least squares solution given by solution (in d) of:

$$\begin{array}{ccc} (A^T A) & d = A^T b \\ 2 \times 2 & 2 \times 1 & 2 \times 1 \end{array}$$

$$\begin{array}{ccc} \left[ \begin{array}{cc} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{array} \right] & \left[ \begin{array}{c} u \\ v \end{array} \right] & = - \left[ \begin{array}{c} \sum I_x I_t \\ \sum I_y I_t \end{array} \right] \\ A^T A & & A^T b \end{array}$$

- The summations are over all pixels in the K x K window
- This technique was first proposed by Lukas & Kanade (1981)

# Eigenvectors of $A^T A$

$$A^T A = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

- Recall the Harris corner detector:  $M = A^T A$  is the *second moment matrix*
- The eigenvectors and eigenvalues of  $M$  relate to edge direction and magnitude
  - The eigenvector associated with the larger eigenvalue points in the direction of fastest intensity change
  - The other eigenvector is orthogonal to it

# Errors in Lucas-Kanade

- The motion is large (larger than a pixel)
  - Not-linear: Iterative refinement
  - Local minima: coarse-to-fine estimation
- A point does not move like its neighbors
  - Motion segmentation
- Brightness constancy does not hold
  - Do exhaustive neighborhood search with normalized correlation - tracking features – maybe SIFT – more later....

# Errors in Lucas-Kanade


- The motion is large (larger than a pixel)
  - Not-linear: Iterative refinement
  - Local minima: coarse-to-fine estimation
- A point does not move like its neighbors
  - Motion segmentation
- Brightness constancy does not hold
  - Do exhaustive neighborhood search with normalized correlation - tracking features – maybe SIFT – more later....

# Not tangent: Iterative Refinement

## Iterative Lukas-Kanade Algorithm


1. Estimate velocity at each pixel by solving Lucas-Kanade equations
2. Warp  $I_t$  towards  $I_{t+1}$  using the estimated flow field
  - - *use image warping techniques*
3. Repeat until convergence

# Optical Flow: Iterative Estimation




(using  $d$  for *displacement* here instead of  $u$ )

# Optical Flow: Iterative Estimation






# Optical Flow: Iterative Estimation



# Optical Flow: Iterative Estimation



# Revisiting the small motion assumption




- Is this motion small enough?
  - Probably not—it's much larger than one pixel (2<sup>nd</sup> order terms dominate)
  - How might we solve this problem?

# Optical Flow: Aliasing




Temporal aliasing causes ambiguities in optical flow because images can have many pixels with the same intensity.

I.e., how do we know which ‘correspondence’ is correct?




To overcome aliasing: coarse-to-fine estimation.

# Reduce the resolution!



# Coarse-to-fine optical flow estimation




Gaussian pyramid of image 1

$u=1.25$  pixels

$u=2.5$  pixels

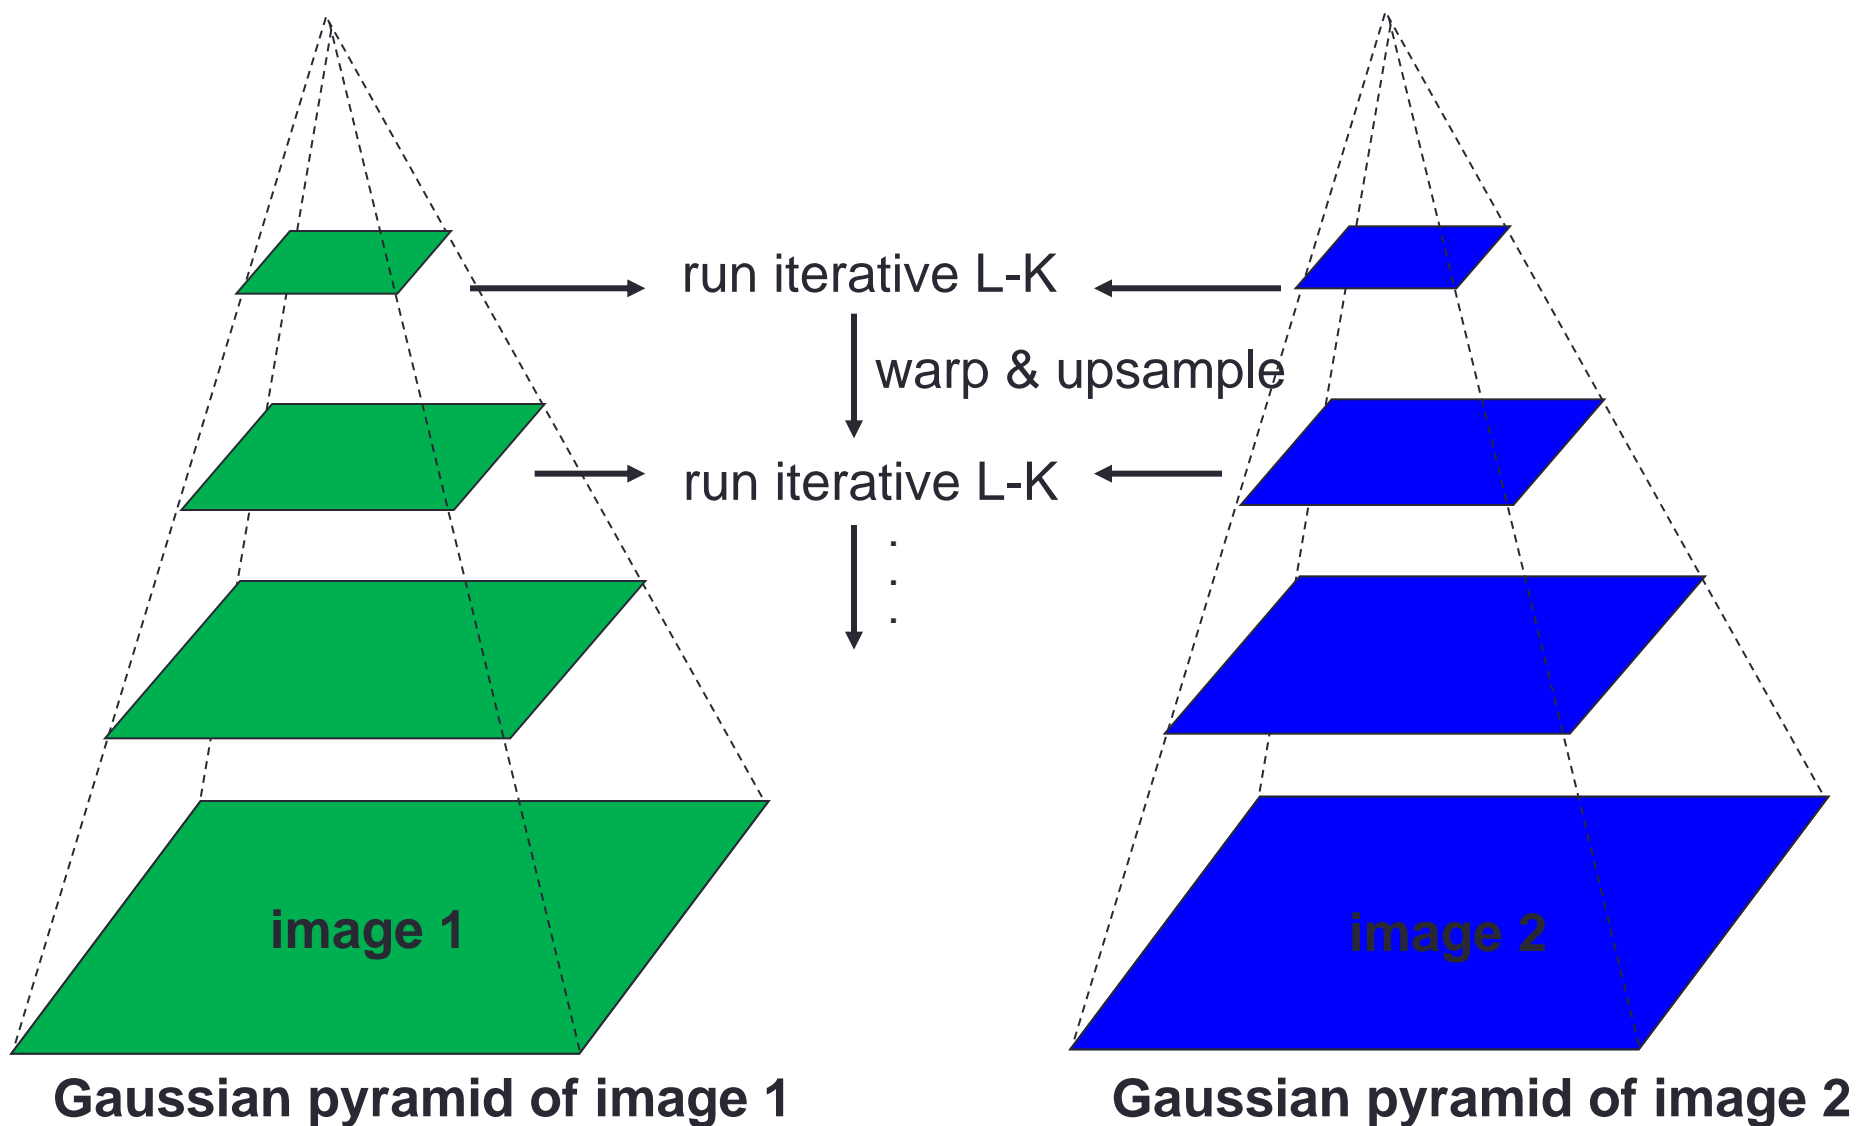
$u=5$  pixels

$u=10$  pixels



Gaussian pyramid of image 2


# Coarse-to-fine optical flow estimation



# Multi-scale



# Remember: Image sub-sampling




1/4



1/8

Throw away every other row and column to create a  $1/2$  size image  
- called *image sub-sampling*

# Bad image sub-sampling



1/2



1/4 (2x zoom)




1/8 (4x zoom)

Aliasing! What do we do?



# Gaussian (lowpass) pre-filtering



Gaussian  $1/2$




G  $1/4$



G  $1/8$

Solution: filter the image, *then* subsample

# Subsampling with Gaussian pre-filtering



Gaussian  $1/2$



G  $1/4$



G  $1/8$

# Band-pass filtering


Gaussian Pyramid (low-pass images)



Laplacian Pyramid (subband images)

*These are “bandpass” images (almost).*

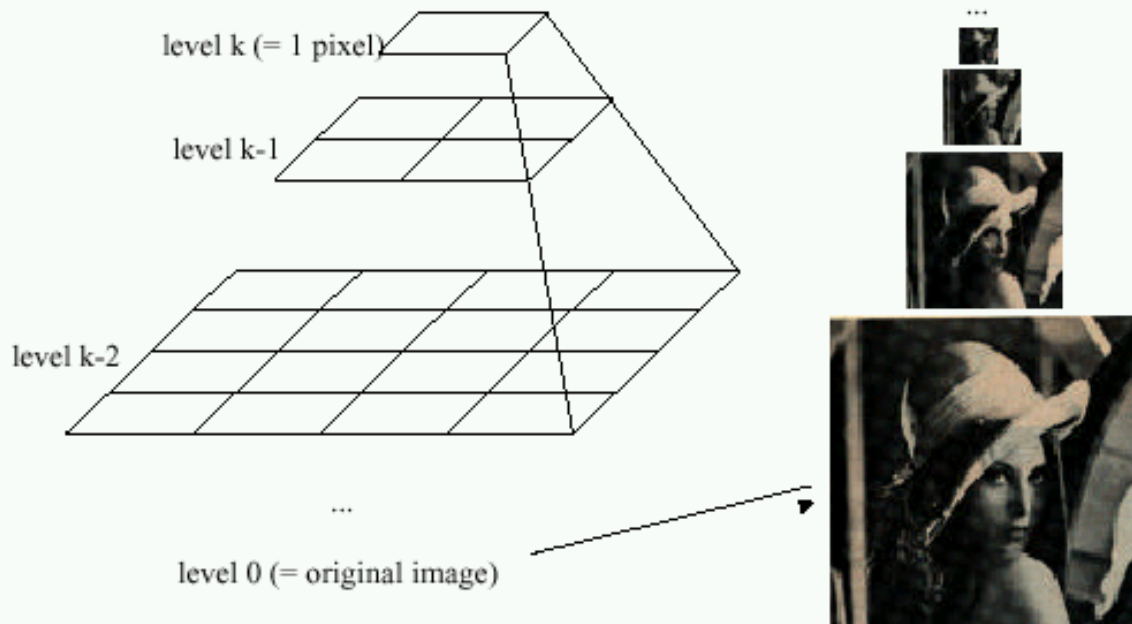
# Laplacian Pyramid



- How can we reconstruct (collapse) this pyramid into the original image?

# Image Pyramids

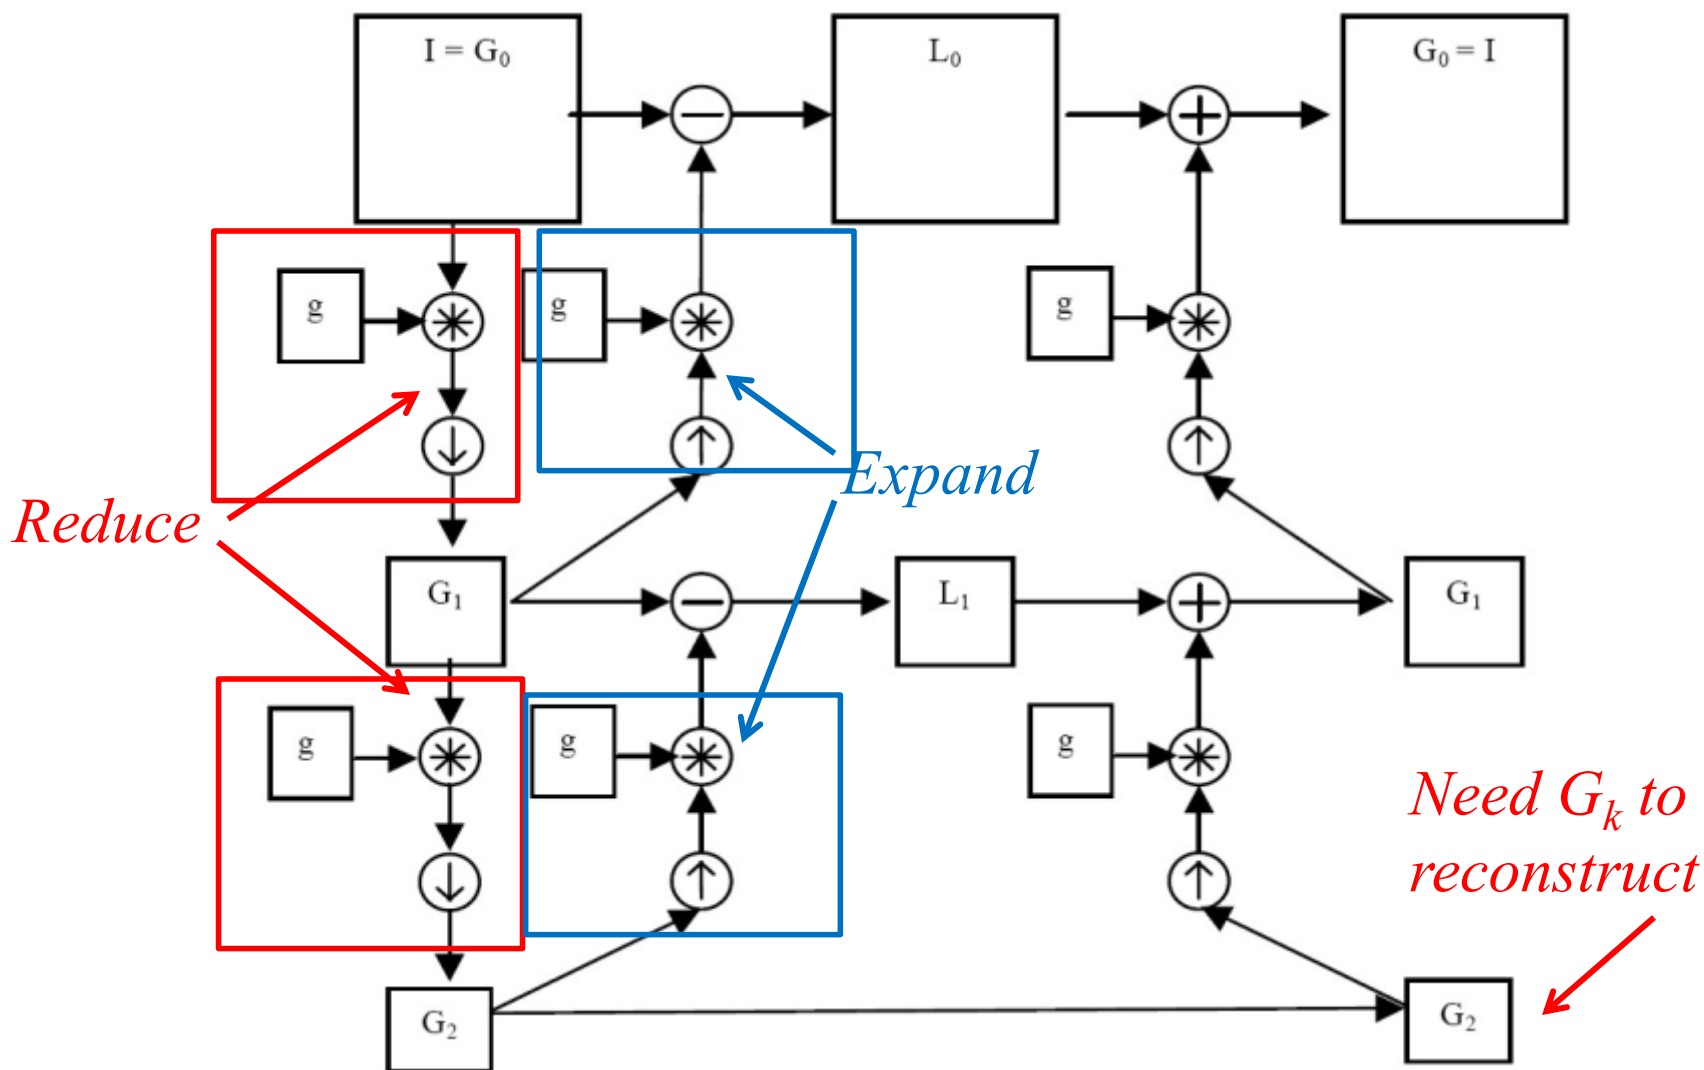
Idea: Represent  $N \times N$  image as a “pyramid” of  $1 \times 1, 2 \times 2, 4 \times 4, \dots, 2^k \times 2^k$  images (assuming  $N = 2^k$ )



Known as a **Gaussian Pyramid** [Burt and Adelson, 1983]


- In computer graphics, a *mip map* [Williams, 1983]

# Computing the Laplacian Pyramid






# Reduce and Expand



(a)

*Reduce*



(b)

*Expand*

# What can you do with band limited imaged?



# Apples and Oranges in bandpass

 $L_0$ 


(a)




(b)




(c)

 $L_2$ 

(d)



(e)



(f)

 $L_4$ 

(g)



(h)



(i)

Reconstructed



(j)



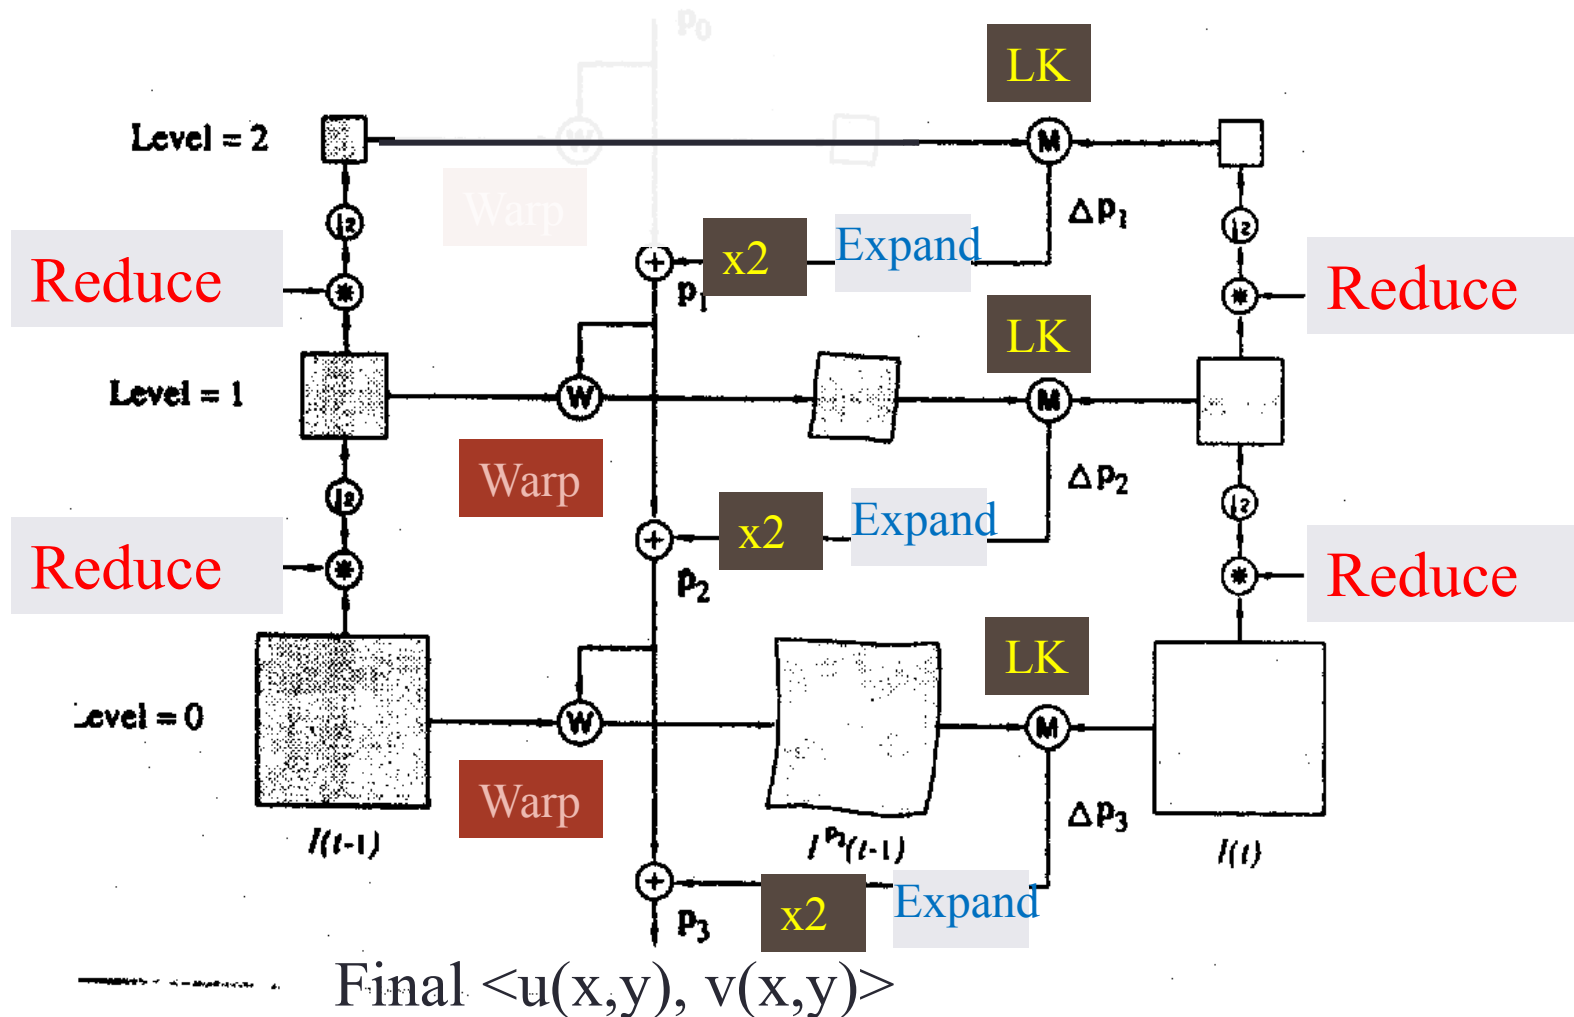
(k)



(l)

# Applying pyramids to LK

# Coarse-to-fine global motion estimation




# Multi-resolution Lucas Kanade Algorithm

Compute Iterative LK at highest level


For Each Level  $i$

- Take flow  $u(i-1)$ ,  $v(i-1)$  from level  $i-1$
- Upsample the flow to create  $u^*(i)$ ,  $v^*(i)$  matrices of twice resolution for level  $i$ .
- Multiply  $u^*(i)$ ,  $v^*(i)$  by 2
- Compute  $I_t$  from a block displaced by  $u^*(i)$ ,  $v^*(i)$
- Apply LK to get  $u'(i)$ ,  $v'(i)$  (the correction in flow)
- Add corrections  $u'(i)$ ,  $v'(i)$  to obtain the flow  $u(i)$ ,  $v(i)$  at  $i^{\text{th}}$  level, i.e.,  
 $u(i)=u^*(i)+u'(i)$ ,  $v(i)=v^*(i)+v'(i)$


# Optical Flow Results



# Optical Flow Results



Lucas-Kanade with Pyramids






# Moving to models


- Previous method(s) give dense flow with little or no constraint between locations (smoothness is either explicit or implicit).
- Suppose you “know” that motion is constrained, e.g.
  - Small rotation about horizontal or vertical axis (or both) that is very close to a translation.
  - Distant independent moving objects
- In this case you might “model” the flow...

*Ready for another old slide?*

# Motion models




**Translation**



**2 unknowns**

**Similarity**




**4 unknowns**

**Affine**



**6 unknowns**

**Perspective**



**8 unknowns**

# Focus of Expansion (FOE) - Example

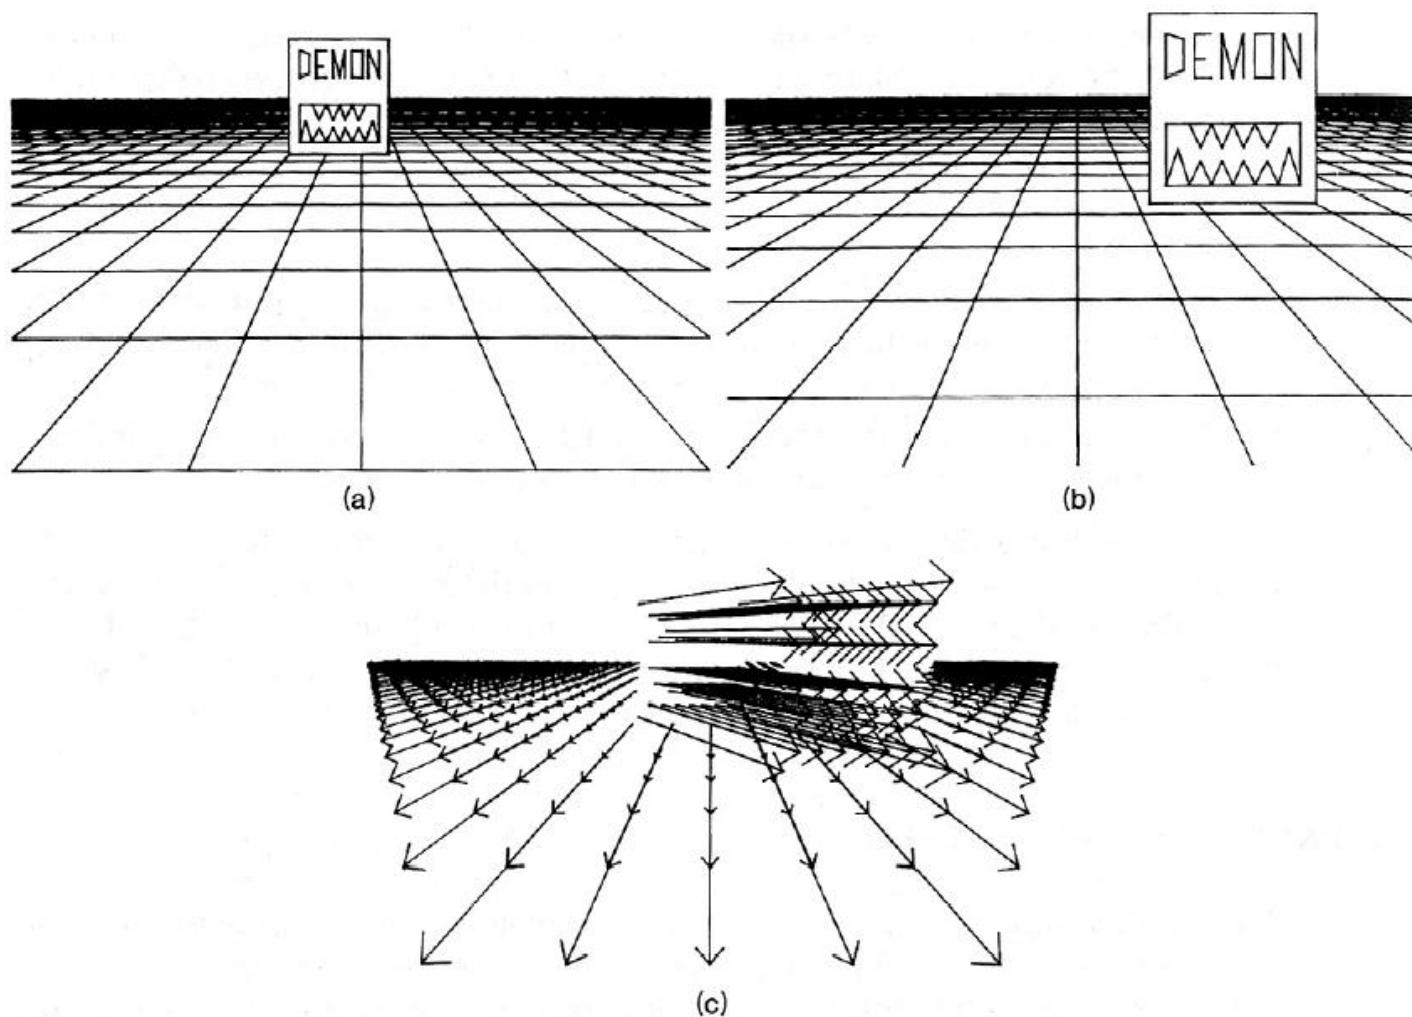


Fig. 7.3 FOE for rectilinear observer motion. (a) An image. (b) Later image. (c) Flow shows different FOEs for static floor and moving object.

# Full motion model

- From physics or elsewhere:

$$\mathbf{V} = \boldsymbol{\Omega} \times \mathbf{R} + \mathbf{T}$$

$$\begin{bmatrix} V_X \\ V_Y \\ V_Z \end{bmatrix} \approx \begin{bmatrix} 0 & -\omega_Z & \omega_Y \\ \omega_Z & 0 & -\omega_X \\ -\omega_Y & \omega_X & 0 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix} + \begin{bmatrix} V_{T_x} \\ V_{T_y} \\ V_{T_z} \end{bmatrix}$$

$$[a_x] = \begin{bmatrix} 0 & -a_3 & a_2 \\ a_3 & 0 & -a_1 \\ -a_2 & a_1 & 0 \end{bmatrix}$$

$$\begin{bmatrix} V_X \\ V_Y \\ V_Z \end{bmatrix} = \text{Velocity Vector}$$

$$\begin{bmatrix} V_{T_x} \\ V_{T_y} \\ V_{T_z} \end{bmatrix} = \text{Translational Component of Velocity}$$

$$\begin{bmatrix} \omega_X \\ \omega_Y \\ \omega_Z \end{bmatrix} = \text{Angular Velocity}$$

# General motion

Take derivatives:

$$x = f \frac{X}{Z}$$

$$y = f \frac{Y}{Z}$$

$$u = v_x = f \frac{ZV_x - XV_z}{Z^2} = f \frac{V_x}{Z} - \left( f \frac{X}{Z} \right) \frac{V_z}{Z} = f \frac{V_x}{Z} - x \frac{V_z}{Z}$$

$$v = v_y = f \frac{ZV_y - YV_z}{Z^2} = f \frac{V_y}{Z} - \left( f \frac{Y}{Z} \right) \frac{V_z}{Z} = f \frac{V_y}{Z} - y \frac{V_z}{Z}$$

$$\begin{bmatrix} u(x, y) \\ v(x, y) \end{bmatrix} = \frac{1}{Z(x, y)} \mathbf{A}(x, y) \mathbf{T} + \mathbf{B}(x, y) \mathbf{\Omega}$$

Why is Z only here?

$$\mathbf{A}(x, y) = \begin{bmatrix} -f & 0 & x \\ 0 & -f & y \end{bmatrix} \quad \mathbf{B}(x, y) = \begin{bmatrix} (xy)/f & -(f + x^2)/f & y \\ (f + y^2)/f & -(xy)/f & -x \end{bmatrix}$$

Where  $\mathbf{T}$  is translation vector,  $\mathbf{\Omega}$  is rotation

# If a plane and perspective...

$$aX + bY + cZ + d = 0$$

$$u(x, y) = a_1 + a_2x + a_3y + a_7x^2 + a_8xy$$

$$v(x, y) = a_4 + a_5x + a_6y + a_7xy + a_8y^2$$

If a plane and orthographic...

$$u(x, y) = a_1 + a_2x + a_3y$$

$$v(x, y) = a_4 + a_5x + a_6y$$

*Affine!*

# Affine motion

$$u(x, y) = a_1 + a_2x + a_3y$$

$$v(x, y) = a_4 + a_5x + a_6y$$


- Substituting into the brightness constancy equation:

$$I_x \cdot u + I_y \cdot v + I_t \approx 0$$

$$I_x(a_1 + a_2x + a_3y) + I_y(a_4 + a_5x + a_6y) + I_t \approx 0$$

- Each pixel provides 1 linear constraint in 6 unknowns
- Least squares minimization:

$$Err(\vec{a}) = \sum \left[ I_x(a_1 + a_2x + a_3y) + I_y(a_4 + a_5x + a_6y) + I_t \right]^2$$





# Affine motion

- Can sum gradients over window or entire image:

$$Err(\vec{a}) = \sum \left[ I_x(a_1 + a_2x + a_3y) + I_y(a_4 + a_5x + a_6y) + I_t \right]^2$$

- Minimize squared error (robustly)

$$\begin{bmatrix} I_x & I_x x_1 & I_x y_1 & I_y & I_y x_1 & I_y y_1 \\ I_x & I_x x_2 & I_x y_2 & I_y & I_y x_2 & I_y y_2 \\ & & \vdots & & & \\ & & \vdots & & & \\ I_x & I_x x_n & I_x y_n & I_y & I_y x_n & I_y y_n \end{bmatrix} \cdot \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \\ a_6 \end{bmatrix} = - \begin{bmatrix} I_t^1 \\ I_t^2 \\ \vdots \\ I_t^n \end{bmatrix}$$

- This is an example of parametric flow – can substitute any linear model easily. Others with some work.

# Hierarchical model-based flow

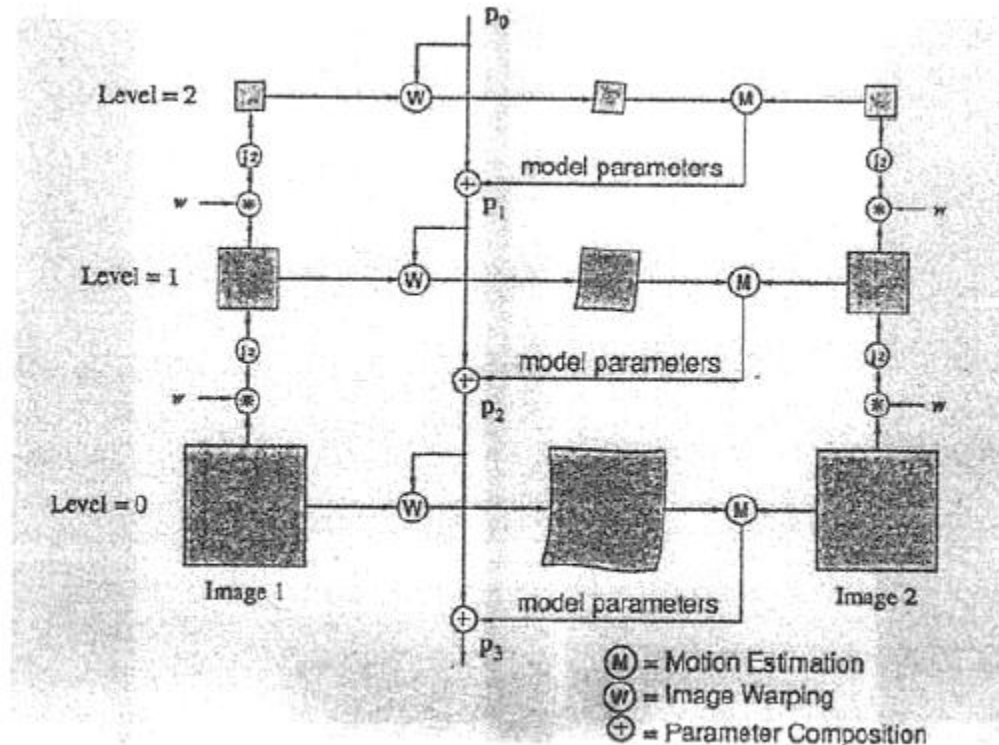



Fig. 1. Diagram of the hierarchical motion estimation framework.

James R. Bergen, P. Anandan, Keith J. Hanna, Rajesh Hingorani:  
 "Hierarchical Model-Based Motion Estimation," ECCV 1992: 237-252

# Now, if different motion regions...




# Layered motion

- Basic idea: break image sequence into “layers” each of which has a coherent motion



# What are layers?

- Each layer is defined by an alpha mask and an affine motion model



# Motion segmentation with an affine model

$$u(x, y) = a_1 + a_2 x + a_3 y$$

$$v(x, y) = a_4 + a_5 x + a_6 y$$

Local flow  
estimates

# Motion segmentation with an affine model

$$u(x, y) = a_1 + a_2 x + a_3 y$$

$$v(x, y) = a_4 + a_5 x + a_6 y$$

Equation of a plane  
(parameters  $a_1, a_2, a_3$  can be  
found by least squares)


# Motion segmentation with an affine model

$$u(x, y) = a_1 + a_2 x + a_3 y$$


$$v(x, y) = a_4 + a_5 x + a_6 y$$

Equation of a plane  
(parameters  $a_1, a_2, a_3$  can be found by least squares)


## 1D example




True flow



Local flow estimate



Segmented estimate



“Foreground”

“Background”

Line fitting






Occlusion



# How do we estimate the layers?

- Compute local flow in a coarse-to-fine fashion
- Obtain a set of initial affine motion hypotheses
  - Divide the image into blocks and estimate affine motion parameters in each block by least squares
    - Eliminate hypotheses with high residual error
  - Perform k-means clustering on affine motion parameters
    - Merge clusters that are close and retain the largest clusters to obtain a smaller set of hypotheses to describe all the motions in the scene
- Iterate until convergence:
  - Assign each pixel to best hypothesis
    - Pixels with high residual error remain unassigned
  - Perform region filtering to enforce spatial constraints
  - Re-estimate affine motions in each region

# Example result



# Recovering image motion

- Direct-methods (e.g. optical flow)
  - Directly recover image motion from spatio-temporal image brightness variations
  - Dense, local motion fields, but more sensitive to appearance variations
  - Suitable for video and when image motion is small ( $< 10$  pixels)
- Feature-based methods (e.g. SIFT, Ransac, regression)
  - Extract visual features (corners, textured areas) and track them - sometimes over multiple frames
  - Sparse motion fields, but possibly robust tracking
    - Good for global motion
  - Suitable especially when image motion is large (10-s of pixels)

# End CS4495